

# Container Platforms

What is available and which should I use?

Lucas Chaufournier

June 7, 2017

# Table of Contents

## General Overview

- Available Platforms

- Docker

- runc

- rkt

- LXC

## Using Containers

## cgroups and Containers

# Available Platforms

- ▶ Docker
- ▶ runc
- ▶ rkt
- ▶ LXC/LXD
- ▶ Openvz



# Docker

- ▶ Most well known platform
- ▶ Common container images are readily available for immediate use (i.e. Docker Hub)
- ▶ Meant to only run one process at a time.
- ▶ Good manageability tools
- ▶ Versioned /layered images
- ▶ Good Marketing and Branding

# runc

- ▶ Most lightweight of all platforms
- ▶ Lots of manual configurations.
- ▶ Similar interface to docker for setting resource limits
- ▶ Relies on docker containers for creating rootfs
- ▶ Export docker rootfs and use json to setup a container.

- ▶ From CoreOS, designed to address some of the security concerns of docker.
- ▶ Use of prepackaged containers from `quay.io` or docker containers
- ▶ Works with kubernetes
- ▶ Uses Systemd for limited container resource management
- ▶ Pods act as a grouping of app images. Talk to each other over localhost. Similar setup to kubernetes pods.
- ▶ Pods mean that all apps in a pod are scheduled together on the same machine and configured with the same resources.

# LXC

- ▶ From the team at ubuntu.
- ▶ Use preconfigured images from image store.
- ▶ Creating new images from scratch not as convenient as other platforms.
- ▶ LXC containers have an init process and can run multiple processes. Meant to be a machine container rather than an app container.
- ▶ Most similar to a VM.

# Creating an Image

## Docker

```
$ vim Dockerfile
{EDIT DOCKERFILE}
$ Docker build .
```

## runc

```
$ mkdir /c1 && cd /c1
$ mkdir rootfs
$ docker export $(docker create busybox) | tar -C rootfs -xvf -
$ runc spec
$ runc create c1
```

## rkt

```
$ acbuild begin
$ acbuild set-name example.com/hello
$ acbuild copy hello /bin/hello
$ acbuild set-exec /bin/hello
$ acbuild write hello-latest-linux-amd64.aci
$ acbuild end
```

## LXC

```
$ lxc-create -t download -n c1
```



# Starting an Image

## Docker

```
$ docker run -it ubuntu /bin/bash
$ docker run -d httpd
```

## runc

```
# cd /c1 && runc run c1
# runc start c1
```

## rkt

```
# rkt run --interactive quay.io/coreos/alpine-sh
# rkt run --interactive docker://ubuntu --insecure-options=image
# systemd-run --slice=machine rkt run quay.io/coreos/alpine-sh
# rkt run example.com/app1 example.com/app2 ##Pod
```

## LXC

```
$ lxc start c1 && lxc exec c1 -- /bin/bash
```

# Stopping a Container

## Docker

```
$ docker stop c4a5ec20a9ec
```

## runc

```
$ runc kill c1 KILL
```

## rkt

```
$ Does not exist
```

## LXC

```
$ lxc stop c1
```

# Deleting a Container

## Docker

```
$ docker rm c4a5ec20a9ec  
$ docker rmi httpd
```

## runc

```
$ runc delete c1
```

## rkt

```
$ rkt rm 203d0797
```

## LXC

```
$ lxc delete c1
```

# Mounting a volume

We are going to mount a host directory `/data` in the container as `/data2`.

## Docker

```
$ docker run -it -v /data:/data2 ubuntu /bin/bash
```

## runc

```
$ Does not exist
```

## rkt:

```
$ rkt run c1 --volume logs,kind=host,source=/data --mount volume=logs,target=/data2
```

## LXC

```
# lxc config device add c1 sdb disk source=/data path=data2
```

# Setting cgroup Limits

We are going to set the maximum memory limits to 8gb. **Note:** You can always use manually cgroups interface for setting limits.

## Docker

```
$ docker run -it --memory=8g ubuntu /bin/bash
```

## runc

```
$ runc update --memory 8000000000 c1
```

## rkt

Use systemd config:

```
[Service]
```

```
Slice=machine.slice
```

```
MemoryLimit=8G
```

```
ExecStart=/usr/bin/rkt run --interactive docker://ubuntu --insecure--options=image --volume
```

```
ExecStopPost=/usr/bin/rkt gc --mark-only
```

```
KillMode=mixed
```

```
Restart=always
```

## LXC

```
$ lxc config set c1 limits.memory 8g
```

# Notes on cgroup interaction

So how does each interface with cgroups?

## Docker and runc

- ▶ Docker runs atop runc.
- ▶ Relies on libcontainer from Docker.
- ▶ libcontainer provides all functionality for creating containers.
- ▶ libcontainer allows interfacing with cgroups using the native interface or systemd.

## rkt

- ▶ Takes advantage of the systemd interface for cgroups.
- ▶ Relies on setting resource limits in a systemd service definition rather than implementing their own system.
- ▶ Only supports: cpu, cpuset and memory.

## LXC

- ▶ Uses liblxc for interacting with cgroups.
- ▶ Filesystem limits can only be set when using zfs or btrfs.